

LAM — Semantic Variations

Neil Ghani and Conor McBride

October 9, 2018

Our grammar of λ -terms is given by

```
<term>
 ::= <var>
    | <term> <term>
    | \ <var> -> <term>
    | <number>
    | (<term> + <term>)
```

If we were real purists, we would omit numbers and addition: we include them to represent ‘ordinary’ computation, where the job of the λ -terms is to route data to whichever computation is appropriate.

1 Small Step Reduction

We write f , s and t for terms, and m and n for numerical constants. Our judgement form is

$$\boxed{s \rightsquigarrow t}$$

with s as an input and t as an output.

The key axioms which do the work are β -reduction and numerical addition.

$$\overline{(\backslash x \rightarrow t[x]) \ s \rightsquigarrow t[s]} \qquad \overline{(m + n) \rightsquigarrow m + n}$$

They are not, however, enough: before you make a reduction step, you have to find a redex! We add *contextual closure* rules:

$$\begin{array}{c} \frac{f \rightsquigarrow f'}{f \ s \rightsquigarrow f' \ s} \qquad \frac{s \rightsquigarrow s'}{f \ s \rightsquigarrow f \ s'} \\[10pt] \frac{t[x] \rightsquigarrow t'[x]}{\backslash x \rightarrow t[x] \rightsquigarrow \backslash x \rightarrow t'[x]} \\[10pt] \frac{s \rightsquigarrow s'}{(s + t) \rightsquigarrow (s' + t)} \qquad \frac{t \rightsquigarrow t'}{(s + t) \rightsquigarrow (s + t')} \end{array}$$

The contextual closure rules are very far from arbitrary: they are generated from the grammar, with one rule for every way to put a subterm inside a term. Their impact is to allow β -reduction and numerical addition *anywhere* inside a term. These rules are highly nondeterministic: when there are multiple redexes, they do not tell you which one to choose.

2 Big Step Reduction

Here is one way to give a big step semantics for LAM. Before we can write the rules, we need to say ‘what success looks like’: to what *values* are we trying to compute the terms?

In our language of functions and numbers, let us say that a function is a value if it is ready to be applied, and a number is a value if it is ready to be added.

```

<value>
  ::= \ <var> -> <term>
      | <number>

```

Note that for functions we make no demands on the body of the abstraction. Let us write u and v for values. Our judgement form is

$$\boxed{t \Downarrow v}$$

with t as input and v as output. We have only four rules. Two tell us how to stop

$$\frac{}{\lambda x. t[x] \Downarrow \lambda x. t[x]} \quad \frac{}{n \Downarrow n}$$

and two tell us how to go.

$$\frac{f \Downarrow \lambda x. t[x] \quad t[s] \Downarrow v}{f s \Downarrow v} \quad \frac{s \Downarrow m \quad t \Downarrow n}{(s + t) \Downarrow m + n}$$

There are several things to notice about this semantics:

- It does not give *every* term a value. (Think about which terms do not get a value, and why. There are several kinds of ‘bad’.)
- It has a deterministic evaluation strategy.
- It never computes inside the body of an abstraction: rather, it leaves the body untouched until the function’s argument turns up.
- It is still reliant on substitution.
- A function’s arguments are not evaluated until after substitution: this is a win if the bound variable is unused, but unfortunate if it is used a lot.

3 Closures and Environments

There is no particularly good reason why values should be a subset of terms. They could be something apart. Here is a way to represent function values more efficiently, getting rid of substitution.

```

<val>
  ::= [<env>] <var> -> <term>
      | <number>

```

```

<env>
  ::=
      | <env>, <var> = <val>

```

Our function values are now given as structures called *closures* which pack up an *environment* along with an abstraction. An environment is a mapping of variables to values. A closure captures the situation where we are not quite ready to evaluate the body of an abstraction: we know the values of all the *free* variables, but not the value of the bound variable.

We write u and v for ‘vals’ and γ for environments. Our judgement form now takes an environment, to explain what the free variables mean.

$$\boxed{\gamma | t \Downarrow v}$$

We expect every free variable to have an entry in the environment. We make the further assumption that all the free variables have different names: this is achievable by α -conversion (but there is a better way).

We have one rule per construct.

$$\frac{\overline{\gamma|\backslash x \rightarrow t \Downarrow [\gamma]x \rightarrow t} \quad \overline{\gamma|n \Downarrow n} \quad \overline{\gamma, x = v, \gamma'|x \Downarrow v}}{\frac{\gamma|f \Downarrow [\gamma']x \rightarrow t \quad \gamma|s \Downarrow u \quad \gamma', x = u|t \Downarrow v}{\gamma|f s \Downarrow v} \quad \frac{\gamma|s \Downarrow m \quad \gamma|t \Downarrow n}{\gamma|(s + t) \Downarrow m + n}}$$

In a real implementation, we get rid of variable *names* entirely. We replace names variable usage sites by numbers called *de Bruijn indices*, locally distinguished from numerical constants by a # marker, which tell us how many λ s to jump over before we find the one which bound the variable. That is

$$\backslash f \rightarrow \backslash x \rightarrow f (f x)$$

becomes

$$\backslash \rightarrow \backslash \rightarrow \#1 (\#1 \#0)$$

A de Bruijn index is just the thing to tell us *where* in γ , now an *array* of values (with 0 the rightmost index), is the value we want.

4 Thunks